

```
In [1]: #1 call
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statistics as stc
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [2]: df=pd.read_csv("googleplaystore.csv")
```

```
In [4]: print("Count of null values in data")
df.isnull().sum()
```

Count of null values in data

```
Out[4]: App          0
Category          0
Rating          1474
Reviews          0
Size             0
Installs         0
Type             1
Price            0
Content Rating    1
Genres           0
Last Updated     0
Current Ver      8
Android Ver      3
dtype: int64
```

```
In [5]: df.dropna(inplace=True)
print("Check for null values after removing nulls")
df.isnull().sum()
```

Check for null values after removing nulls

```
Out[5]: App          0
Category          0
Rating           0
Reviews          0
Size             0
Installs         0
Type             0
Price            0
Content Rating    0
Genres           0
Last Updated     0
Current Ver      0
```

Android Ver 0
dtype: int64

```
In [6]: df=df[-df['Size'].str.contains('Var')]
```

```
In [7]: df.loc[:, 'SizeNum'] =df.Size.str.rstrip('Mk+')  
df.SizeNum=pd.to_numeric(df['SizeNum'])  
df.SizeNum.dtype
```

Out[7]: dtype('float64')

```
In [8]: df['SizeNum']=np.where(df.Size.str.contains('M'),df.SizeNum*1000, df.SizeNum)
```

```
In [9]: # Size no more needed, replace it with SizeNum and drop SizeNum  
df.Size=df.SizeNum  
df.drop('SizeNum',axis=1,inplace=True)  
#df
```

```
In [10]: df.Reviews = pd.to_numeric(df.Reviews)
```

```
In [11]: df.Reviews.dtype
```

Out[11]: dtype('int64')

```
In [12]: df['Installs']=df.Installs.str.replace("+", "")
```

C:\Users\MYPC\AppData\Local\Temp\ipykernel_4732\519759075.py:1: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['Installs']=df.Installs.str.replace("+", "")
```

```
In [13]: df.Installs=df.Installs.str.replace(",", "")  
df.Installs=pd.to_numeric(df.Installs)  
df.Installs.dtype
```

Out[13]: dtype('int64')

```
In [14]: df.Price=df.Price.str.replace("$", "")  
df.Price=pd.to_numeric(df.Price)  
df.Price.dtype
```

C:\Users\MYPC\AppData\Local\Temp\ipykernel_4732\3806789969.py:1: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will **not** be treated as literal strings when regex=True.

```
df.Price=df.Price.str.replace("$","")
dtype('float64')
```

Out[14]:

In [15]:

```
df=df[(df.Rating>=1) & (df.Rating<=5) ]
```

In [16]:

```
len(df.index)
```

Out[16]:

7723

In [17]:

```
df.drop(df.index[df.Reviews>df.Installs],axis=0,inplace=True)
len(df.index)
```

Out[17]:

7717


In [19]:

```
index_free_and_price_gt_0=df.index[((df.Type=='Free')&(df.Price>0))]
if len(index_free_and_price_gt_0)>0:
    print("Dropping following indices:",index_free_and_price_gt_0)
    df.drop(index_free_and_price_gt_0,axis=0,inplace=True)
else:
    print("There is no Free Apps with price >0")
```

There is no Free Apps with price >0

In [20]:

```
ax = sns.boxplot(x='Price', data=df)
```



In [22]:

```
price_std=stc.stdev(df.Price)
```

```
price_std
```

```
Out[22]: 17.414783874309933
```

```
In [23]: price_mean=stc.mean(df.Price)
price_mean
```

```
Out[23]: 1.128724893093171
```

```
In [24]: price_outlier_uplimit=price_mean+3*price_std
price_outlier_uplimit
```


```
Out[24]: 53.37307651602297
```

```
In [25]: #price_outlier_downlimit=price_mean-3*price_std
#price_outlier_downlimit
#df[df.Price>price_outlier_uplimit]
print("# of upper outliers is ",len(df[(df.Price>price_outlier_uplimit) ]))

# of upper outliers is 17
```

```
In [26]: #df[df.Price<price_outlier_downlimit]
#print("# of lower outliers is ",len(df[df.Price<price_outlier_downlimit]))
sns.boxplot(x='Reviews',data=df)
```

```
Out[26]: <AxesSubplot:xlabel='Reviews'>
```



```
In [27]: rev_std=stc.stdev(df.Reviews)
rev_std
```

```
1864639.6094670836
```

Out[27]:

```
In [29]: rev_mean=stc.mean(df.Reviews)
rev_mean
```

Out[29]: 295127.5482700531

```
In [30]: rev_outlier_uplimit=rev_mean+3*rev_std
rev_outlier_uplimit
```

Out[30]: 5889046.376671304

```
In [31]: rev_outlier_downlimit=rev_mean-3*rev_std
rev_outlier_downlimit
```

Out[31]: -5298791.280131198


```
In [32]: #df[df.Reviews>rev_outlier_uplimit]
print("# of upper outliers is ",len(df[(df.Reviews>rev_outlier_uplimit) ]))

# of upper outliers is 89
```

```
In [33]: # Since reviews cannot be less than 1, no need to check lower outliers
# remove outliers
#df.drop(df.index[(df.Reviews>rev_outlier_uplimit) ],inplace=True)
#len(df.index)
```


```
In [34]: #sns.boxplot(x='Rating',data=df)
sns.histplot(x='Rating',data=df)
#rating_std=stc.stdev(df.Rating)
#rating_std
#rating_mean=stc.mean(df.Rating)
#rating_mean
#rating_outlier_uplimit=rating_mean+3*rating_std
#rating_outlier_uplimit
#rating_outlier_downlimit=rating_mean-3*rating_std
#rating_outlier_downlimit
# Since max possible value of rating (5) is less than upper limit, no need to
#df[df.Rating<rating_outlier_downlimit]
#print("# of lower outliers is ",len(df[(df.Rating<rating_outlier_downlimit)
#df.drop(df.index[(df.Rating<rating_outlier_downlimit)],inplace=True)
#len(df.index)
```

Out[34]: <AxesSubplot:xlabel='Rating', ylabel='Count'>



```
In [35]: # use log scale to make histogram more representable
sns.histplot(x='Size',data=df,log_scale=True)
```

Out[35]: <AxesSubplot:xlabel='Size', ylabel='Count'>



```
In [36]: df[df.Price>=200]
```

Out[36]:

	App	Category	Rating	Reviews	Size	Installs	Type	Price	Content Rating
4197	most expensive app (H)	FAMILY	4.3	6	1500.0	100	Paid	399.99	Everyone
4362	💎 I'm rich	LIFESTYLE	3.8	718	26000.0	10000	Paid	399.99	Everyone

	App	Category	Rating	Reviews	Size	Installs	Type	Price	Content Rating
4367	I'm Rich - Trump Edition	LIFESTYLE	3.6	275	7300.0	10000	Paid	400.00	Everyone
5351	I am rich	LIFESTYLE	3.8	3547	1800.0	100000	Paid	399.99	Everyone
5354	I am Rich Plus	FAMILY	4.0	856	8700.0	10000	Paid	399.99	Everyone
5355	I am rich VIP	LIFESTYLE	3.8	411	2600.0	10000	Paid	299.99	Everyone
5356	I Am Rich Premium	FINANCE	4.1	1867	4700.0	50000	Paid	399.99	Everyone
5357	I am extremely Rich	LIFESTYLE	2.9	41	2900.0	1000	Paid	379.99	Everyone
5358	I am Rich!	FINANCE	3.8	93	22000.0	1000	Paid	399.99	Everyone
5359	I am rich(premium)	FINANCE	3.5	472	965.0	5000	Paid	399.99	Everyone
5362	I Am Rich Pro	FAMILY	4.4	201	2700.0	5000	Paid	399.99	Everyone
5364	I am rich (Most expensive app)	FINANCE	4.1	129	2700.0	1000	Paid	399.99	Teen
5366	I Am Rich	FAMILY	3.6	217	4900.0	10000	Paid	389.99	Everyone
5369	I am Rich	FINANCE	4.3	180	3800.0	5000	Paid	399.99	Everyone
5373	I AM RICH PRO PLUS	FINANCE	4.0	36	41000.0	1000	Paid	399.99	Everyone



In [37]: `print("# of Apps with price >= 200 = ",len(df[(df.Price>=200)]))`

of Apps with price >= 200 = 15

In [38]: `df.drop(df.index[(df.Price>=200)], inplace=True)
len(df.index)`

Out[38]: 7702

```
In [39]: df.drop(df.index[(df.Reviews>=2000000)], inplace=True)  
len(df.index)
```

Out[39]: 7483

```
In [40]: install_10_perc=np.percentile(df.Installs, 10)  
install_10_perc
```

Out[40]: 1000.0

```
In [41]: install_25_perc=np.percentile(df.Installs, 25)  
install_25_perc
```

Out[41]: 10000.0

```
In [42]: install_50_perc=np.percentile(df.Installs, 50)  
install_50_perc
```

Out[42]: 100000.0

```
In [43]: install_70_perc=np.percentile(df.Installs, 70)  
install_70_perc
```

Out[43]: 1000000.0

```
In [44]: install_90_perc=np.percentile(df.Installs,90)  
install_90_perc
```

Out[44]: 10000000.0

```
In [45]: install_95_perc=np.percentile(df.Installs,95)  
install_95_perc
```


Out[45]: 10000000.0

```
In [46]: install_99_perc=np.percentile(df.Installs,99)  
install_99_perc
```

Out[46]: 50000000.0

```
In [47]: sns.histplot(data=df,x='Installs',log_scale=True)
```


Out[47]: <AxesSubplot:xlabel='Installs', ylabel='Count'>



In [48]: `print("As result, ",len(df[df.Installs >= install_99_perc])," will be dropped")`


As result, 176 will be dropped

In [49]: `df.drop(df.index[df.Installs >= install_99_perc],inplace=True)`
`len(df.index)`

Out[49]: 7307


In [50]: `sns.jointplot(data=df,y='Rating',x='Price')`

Out[50]: <seaborn.axisgrid.JointGrid at 0x1b87a920ca0>




```
In [51]: sns.jointplot(data=df,y='Rating',x='Size')
```

```
Out[51]: <seaborn.axisgrid.JointGrid at 0x1b87aad5640>
```



```
In [52]: sns.jointplot(data=df,y='Rating',x='Reviews')
```

```
Out[52]: <seaborn.axisgrid.JointGrid at 0x1b87c197160>
```




In [53]: `df['Content Rating'].unique()`

Out[53]: `array(['Everyone', 'Teen', 'Everyone 10+', 'Mature 17+',
 'Adults only 18+', 'Unrated'], dtype=object)`


In [54]: `sns.boxplot(data=df, x='Rating', y='Content Rating')`

Out[54]: `<AxesSubplot:xlabel='Rating', ylabel='Content Rating'>`



```
In [55]: a4_dims = (11.7, 10.27)
fig, ax = plt.subplots(figsize=a4_dims)
sns.boxplot(data=df, x='Rating', y='Category', ax=ax)
```

```
Out[55]: <AxesSubplot:xlabel='Rating', ylabel='Category'>
```



```
In [56]: #8.1
inp1=df.copy()
inp1.Reviews=inp1.Reviews.apply(np.log1p)
```

```
In [57]: inp1.Installs=inp1.Installs.apply(np.log1p)
```

```
In [60]: #8.2
inp1.drop(columns=['App', 'Last Updated', 'Current Ver', 'Android Ver'], inplace=
```

```
In [61]: inp1.shape
```

```
Out[61]: (7307, 9)
```

```
In [62]: #8.3
```

```
inp2= pd.get_dummies(inp1)
```

```
In [63]: inp2.shape
```

```
Out[63]: (7307, 158)
```

```
In [64]: data = inp2.drop(columns='Rating')  
data.shape
```

```
Out[64]: (7307, 157)
```

```
In [65]: target = pd.DataFrame(inp2.Rating)  
target.shape
```

```
Out[65]: (7307, 1)
```

```
In [68]: x_train, x_test, y_train, y_test = train_test_split(data, target, test_size=0.2)  
print("x_train shape is ", x_train.shape)  
print("y_train shape is ", y_train.shape)  
print("x_test shape is ", x_test.shape)  
print("y_test shape is ", y_test.shape)
```

```
x_train shape is (5114, 157)  
y_train shape is (5114, 1)  
x_test shape is (2193, 157)  
y_test shape is (2193, 1)
```

```
In [69]: model=LinearRegression()  
model.fit(x_train, y_train)
```

```
Out[69]: LinearRegression()
```

```
In [70]: train_pred=model.predict(x_train)
```

```
In [71]: print("R2 value of the model(by train) is ", r2_score(y_train, train_pred))
```

```
R2 value of the model(by train) is  0.15264772134593874
```

```
In [72]: test_pred=model.predict(x_test)
```

```
In [73]: print("R2 value of the model(by test) is ", r2_score(y_test, test_pred))
```

```
R2 value of the model(by test) is  0.14262263030973144
```

